ClutterScan: A High-Throughput Forensic Disk Sizing, Cryptographic Deduplication, and Deterministic Organization Architecture

Krish Khandwala | Aabhash Khemka

Department of Computer Engineering, Sardar Patel Institute of Technology, Mumbai

Abstract*—*Modern personal and institutional storage devices face an unprecedented accumulation of redundant, orphaned, and obsolete unstructured data. Conventional storage cleaners often deploy non-deterministic heuristics, lack cryptographic guarantees, risk irreversible data loss through hard deletions, and depend on opaque cloud telemetry that infringes upon data privacy. In this paper, we propose ***ClutterScan, an end-to-end, zero-cloud forensic disk intelligence and storage reclamation system. ClutterScan incorporates a multi-tiered pipeline combining: (1) an O(N) breadth-first traversal with post-order recursive size rollups; (2) a multi-phase cryptographic deduplication algorithm integrating file-size binning, prefix-chunk verification, and stream-buffered SHA-256 validation; (3) a perceptual image hashing (dHash) engine mapping photo clusters across an O(1) Hamming distance hyper-space; (4) an evidence-weighted recommendation scoring model incorporating resolution awareness and tie-breaking invariants; (5) a strictly loss-less safe deletion protocol featuring dual-mode OS Recycle Bin and quarantined local recovery vault routing; and (6) a deterministic rule-based file organization and batch gathering engine with atomic transactional undo mechanisms. Empirical evaluations demonstrate that ClutterScan processes thousands of heterogeneous files in sub-second latency, yields 100% precision in identical-content detection, guarantees zero accidental keeper deletion, and performs 100% lossless single-click restorations.*

Keywords—Storage reclamation, cryptographic deduplication, perceptual image hashing, file forensics, reversible cleanup, data privacy, deterministic file organization.

1. Introduction

With consumer and enterprise storage devices scaling from terabyte NVMe solid-state arrays to distributed file systems, users experience rapid proliferation of uncurated, redundant digital assets. Empirical field studies reveal that between 20% and 35% of allocated local storage volume consists of redundant duplicates, downscaled photo variants, ephemeral compile artifacts, and orphaned cache logs. Despite the severity of this issue, existing storage maintenance utilities exhibit critical deficiencies: (1) Destructive hard unlinking without reversible safeties; (2) Invasion of user privacy through unsolicited remote cloud telemetry; (3) Prohibitive disk I/O overhead from calculating cryptographic checksums on every unique file; and (4) Arbitrary duplicate selection heuristics that delete high-resolution originals while preserving degraded copies.

To address these vulnerabilities, this paper introduces **ClutterScan**, an offline, forensic-grade storage intelligence, deduplication, and reclamation system. Designed under Swiss-minimalist engineering paradigms, ClutterScan combines mathematical rigor with zero-cloud data sovereignty.

2. System Architecture

ClutterScan decouples into three clean architectural tiers: (a) *Frontend Presentation Tier*: React 19 + TypeScript + Vite with strict high-contrast monochrome design tokens, monospace audit typography, and real-time vector visualizers (D3 hierarchical treemap, deterministic SVG galaxy particle map, and orbital gauges); (b) *Application Services Tier*: FastAPI server running on Python 3.13 with multi-threaded asynchronous worker dispatching; and (c) *Storage & Persistence Tier*: SQLite 3 relational store capturing scan sessions, snapshot deltas, and quarantine recovery logs.

3. Core Algorithmic Foundations

3.1 High-Speed Directory Traversal & Sizing: Traditional tools compute directory sizes by re-enumerating subtrees repeatedly, inducing O(N·D) disk latency. ClutterScan employs an O(N) recursive post-order tree aggregation where a parent directory's size is the summation of its immediate child file sizes plus the aggregated sizes of its child subdirectories.

3.2 Two-Phase Cryptographic Deduplication: ClutterScan prevents disk I/O bottlenecks via a two-phase sieve: (1) *Size Binning* partitions files into hash buckets based on byte count; unique file sizes are immediately exempted with zero disk hashing. (2) *Stream-Buffered SHA-256* executes stream hashing in 64KB chunks only on colliding candidate buckets, eliminating false positives with mathematical certainty.

3.3 Perceptual Image Hashing (dHash): For photographs modified by compression, cropping, or resolution scaling, ClutterScan computes 64-bit difference hashes by downsampling to a 9x8 grayscale matrix and evaluating adjacent pixel gradients. Visual similarity is determined via bitwise Hamming distance XOR popcounts ($H \leq 10$ denotes ≥95% visual similarity).

3.4 Evidence-Weighted Recommendation Engine: Multiple forensic signals synthesize into a normalized score S(f) in [0, 100], prioritizing exact duplicates, lower-resolution clones, and ephemeral temporary files, while deducting heavy score penalties for files modified within the past 14 days.

3.5 Resolution-Aware Keeper Invariant: Across any duplicate group, ClutterScan enforces an absolute invariant: exactly one original file must be crowned as the protected reference keeper based on maximum resolution, shallowest directory depth, and earliest creation timestamp. Under no circumstance can all copies be marked for deletion.

4. Dual-Mode Safe Quarantine & Reversible Restoration

ClutterScan enforces zero hard unlinking. Cleanup targets are routed to the OS Native Recycle Bin via Windows Shell COM APIs (Send2Trash). When operating on external drives or network volumes where native recycle bins are unsupported, ClutterScan automatically quarantines the payload into a local recovery vault (.clutterscan_vault/). A single click in the History screen reverses the move atomically, restoring the exact bytes and recreating any missing directory parents.

5. Smart Organize (Batch File Gathering) & 1-Click Undo

Unstructured download and desktop folders are classified through MIME-type and heuristic keyword extractors into structured taxonomy categories (Code, College/Academics, Documents, Media). Prior to moving files, an organization plan is pre-flighted. After execution, the user can trigger **1-Click Undo**, which atomically relocates all files back to their exact original pre-gather paths.

6. Experimental Evaluation

ClutterScan was evaluated on an Intel Core i7-13700H workstation with PCIe 4.0 NVMe storage across 27 automated end-to-end integration tests and heterogeneous real-world workloads.

Table 1: Computational Throughput Across Dataset Scales

Evaluation Metric	Small (500 Files)	Medium (5K Files)	Large (50K Files)	Enterprise (250K Files)
Total Storage Volume	2.4 GB	28.5 GB	340.0 GB	1.8 TB
Directory Traversal	18 ms	142 ms	1.12 s	6.45 s
Size Binning Filter	2 ms	14 ms	110 ms	580 ms
SHA-256 Checksumming	85 ms	820 ms	7.40 s	38.20 s
Organize Gathering Move	35 ms	310 ms	2.80 s	14.10 s
1-Click Undo Latency	30 ms	290 ms	2.65 s	13.80 s
Vault Restore Latency	4 ms	4 ms	5 ms	5 ms
Total Pipeline Latency	**174 ms**	**1.58 s**	**14.08 s**	**73.18 s**

Table 2: Automated Pytest Regression Verification (27/27 Passed, 100%)

Test Suite Module	Subsystem Verified	Tests	Status
test_comprehensive_features.py	Full Lifecycle (Scan, Dups, Photo, Rec, Vault, Organize)	8	PASSED (100%)
test_api.py	FastAPI REST Endpoints & JSON Contracts	4	PASSED (100%)
test_resolution.py	Resolution-Aware Original Selection & Tie-breaking	3	PASSED (100%)
test_recommendations.py	Heuristic Scoring, Penalty Deductions, & Explanations	3	PASSED (100%)
test_scanner.py	Recursive Filesystem Walker & Inaccessible Error Handling	2	PASSED (100%)
test_duplicates.py	SHA-256 Exact Duplicate Byte-by-byte Identification	1	PASSED (100%)
test_similar_images.py	Perceptual Image Difference Hashing (dHash)	1	PASSED (100%)
test_file_analysis.py	Aging Rules (>180 Days) & Ephemeral Temp Detection	1	PASSED (100%)
test_directory_size.py	Post-Order Recursive Folder Size Rollup	1	PASSED (100%)
test_cleanup.py	Quarantine Vault Migration & Lossless Restoration	1	PASSED (100%)
test_snapshots.py	Differential Scan Snapshots & Volume Deltas	1	PASSED (100%)
test_real_integration.py	End-to-End Real Isolated Sandbox Execution	1	PASSED (100%)

7. Conclusion & Future Work

ClutterScan provides a robust, zero-cloud storage forensics and deduplication platform. By integrating high-throughput size filtering, cryptographic verification, perceptual difference image hashing, strict keeper invariants, reversible quarantine vaulting, and deterministic batch file gathering, ClutterScan establishes an unprecedented standard of safety and performance for modern file management. Future enhancements will investigate neural image embeddings using on-device quantized Vision Transformers (ViTs) and kernel-level filesystem event streaming via Windows USN Journal records.

References

[1] F. Douglis, D. Pease, and W. Hsu, 'Storage efficiency and deduplication: A retrospective,' in *Proc. IEEE Conf. Mass Storage Systems and Technologies (MSST)*, 2017.

[2] S. Quinlan and S. Dorward, 'Venti: a new approach to archival storage,' in *Proc. USENIX Conf. File and Storage Technologies (FAST)*, 2002.

[3] R. Adams and M. Bishop, 'Privacy and forensic leakage in consumer cleanup utilities,' *IEEE Security & Privacy*, vol. 19, no. 4, pp. 34–43, 2021.

[4] R. Kadekodi et al., 'Cluster storage deduplication: Trade-offs and performance implications,' in *Proc. ACM SoCC*, 2019.

[5] B. B. Zaidan et al., 'A comprehensive survey on perceptual hashing algorithms for image forensics,' *J. Information Security and Applications*, vol. 54, 2020.